

Nginx shows 2 different statuses in the upstream response log

Asked 3 years, 7 months ago Modified 3 years, 7 months ago Viewed 1k times



1



I am running an nginx-ingress controller in a kubernetes cluster and one of my log statements for the request looks like this:

```
upstream_response_length: 0, 840
upstream_response_time: 60.000, 0.760
upstream_status: 504, 200
```

I cannot quite understand what does that mean? Nginx has a response timeout equal to 60 seconds, and tries to request one more time after that (successfully) and logs both requests?

P.S. Config for log format:

```
log-format-upstream: >-
{
  ...
  "upstream_status": "$upstream_status",
  "upstream_response_length": "$upstream_response_length",
  "upstream_response_time": "$upstream_response_time",
  ...
}
```

nginx

kubernetes

nginx-ingress

nginx-log

Share Improve this question

edited Sep 26, 2019 at 10:18

Follow

asked Sep 26, 2019 at 7:33



Ilya Chernomordik

27k ● 26 ● 114 ● 202

Would you please show your configs? – [Yasen](#) Sep 26, 2019 at 10:02

1 Answer

Sorted by:

Highest score (default)



2



According to [split_upstream_var](#) method of [ingress-nginx](#), it [splits](#) results of `nginx` [health checks](#).

Since `nginx` can have [several upstreams](#), your log could be interpreted this way:

1. First upstream is dead (504)

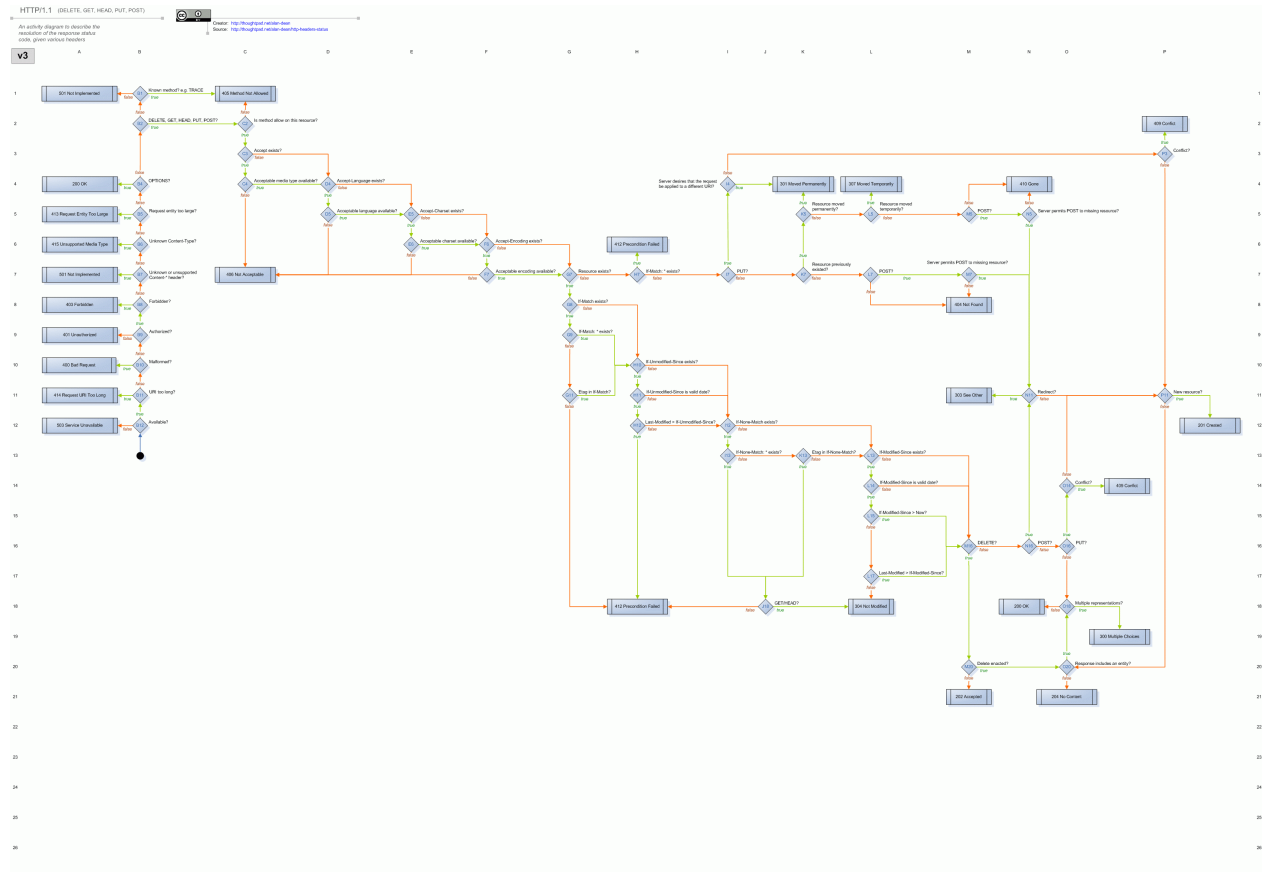
```
upstream_response_length: 0 // response from dead upstream has zero length
upstream_response_time: 60.000 // nginx dropped connection after 60sec
```

```
upstream_status: 504 // response code, upstream doesn't answer
```

2. Second upstream works (200)

```
upstream_response_length: 840 // healthy upstream returned 840b
upstream_response_time: 0.760 // healthy upstream responded in 0.760
upstream_status: 200 // response code, upstream is ok
```

P.S. JFYI, here's a cool HTTP headers state diagram



Share Improve this answer Follow

edited Sep 26, 2019 at 9:27

answered Sep 26, 2019 at 9:22



Yasen

4,135 ● 1 ● 16 ● 25

I guess it's really just one upstream for nginx ingress controller, which is a service in kubernetes. Is it possible that nginx just does another request to upstream when the first is dropped after 60 seconds?


– Ilya Chernomordik Sep 26, 2019 at 9:40



Explore



▼ A (main-loki)



Query patterns ▼ Explain ☐

Builder **Beta** Code

Log browser >

```
max_over_time({namespace="fed-mdecss-nginx-ingress-controller"} |= "GET" |= "v39/mdecss" | json | __error__!="JSONParserErr" |
line_format "{.log}") | json | __error__!="JSONParserErr" | unwrap upstream_response_time [1m]) by (upstream_status)
```

> Options Type: Range



pipeline error: 'SampleExtractionErr' for series: '{__error__="SampleExtractionErr", app="nginx-ingress-controller", component="controller", container="controller", filename="/var/log/pods/fed-mdecss-nginx-ingress-controller_fed-mdecss-nginx-ingress-controller-7b7cd9bbf-qg9cc_26e3c20c-c1a1-4897-9bc7-f7eef55f9990/controller/0.log", instance="fed-mdecss-nginx-ingress-controller", job="fed-mdecss-nginx-ingress-controller/nginx-ingress-controller", log="{\"time_local\": \"10/May/2023:01:31:25 +0000\", \"geoip_city\": \"Jinan\", \"geoip_country_code\": \"CN\", \"remote_addr\": \"58.56.98.14\", \"remote_port\": \"7386\", \"http_token\": \"IO-K9LV-CNNP4-J3ZS6-3TQ7Q-WD8RB-CK25B\", \"http_product_id\": \"\", \"http_pid\": \"ESMv8\", \"host\": \"csp.spn.asiainfo-sec.com\", \"scheme\": \"https\", \"request_uri\": \"/v39/mdecss?version=315\", \"args\": \"version=315\", \"sent_http_connection\": \"close\", \"status\": \"200\", \"request_time\": \"4.274\", \"upstream_response_time\": \"0.075, 4.198\", \"upstream_status\": \"502, 200\", \"server_protocol\": \"HTTP/1.1\", \"request_method\": \"GET\", \"body_bytes_sent\": \"17970578\", \"http_user_agent\": \"\", \"upstream_addr\": \"192.168.20.51:80, 192.168.28.170:80\"}\\n\", namespace=\"fed-mdecss-nginx-ingress-controller\", node_name=\"ip-192-168-21-27.cn-northwest-1.compute.internal\", pod=\"fed-mdecss-nginx-ingress-controller-7b7cd9bbf-qg9cc\", upstream_response_time=\"0.075, 4.198\", upstream_status=\"502, 200\"}'. Use a label filter to intentionally skip this error. (e.g | __error__!="SampleExtractionErr"). To skip all potential errors you can match empty errors.(e.g __error__="") The label filter can also be specified after unwrap. (e.g | unwrap latency | __error__="")

+ Add query

🕒 Query history

🔍 Inspector